```python
# This Python 3 environment comes with many helpful analytics
libraries installed
# It is defined by the kaggle/python Docker image:
https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
import matplotlib.pyplot as plt
from sklearn.metrics import precision_score, recall_score, f1_score,
accuracy_score
import torch
import torch.nn as nn
import torch.optim as optim
import torchvision.transforms as transforms
import torchvision.datasets as datasets
from torchvision.models import vgg16


# Input data files are available in the read-only "../input/"
directory
# For example, running this (by clicking run or pressing Shift+Enter)
will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
that gets preserved as output when you create a version using "Save &
Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
be saved outside of the current session

# Set the path to the root directory of your dataset
train_data_path = '/kaggle/input/neu-metal-surface-defects-data/NEU
Metal Surface Defects Data/train'
test_data_path = '/kaggle/input/neu-metal-surface-defects-data/NEU
Metal Surface Defects Data/test'
valid_data_path = '/kaggle/input/neu-metal-surface-defects-data/NEU
Metal Surface Defects Data/valid'

# Data transformations
data_transforms = transforms.Compose([
    transforms.Resize((224, 224)),  # Resize the images to 224x224
(required for ResNet-34)
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229,
0.224, 0.225])  # Normalize the images
```

```python
])

train_dataset = datasets.ImageFolder(train_data_path,
transform=data_transforms)
test_dataset = datasets.ImageFolder(test_data_path,
transform=data_transforms)
valid_dataset = datasets.ImageFolder(valid_data_path,
transform=data_transforms)

train_loader = torch.utils.data.DataLoader(train_dataset,
batch_size=64, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_dataset, batch_size=64,
shuffle=False)
valid_loader = torch.utils.data.DataLoader(valid_dataset,
batch_size=64, shuffle=False)

# Function to plot a grid of images
def imshow_grid(images, labels, classes, rows=4, cols=4):
    fig, axes = plt.subplots(rows, cols, figsize=(10, 10))

    for i, ax in enumerate(axes.flat):
        img = images[i]
        label = labels[i]
        ax.imshow(np.transpose(img, (1, 2, 0)))
        ax.set_title(classes[label])
        ax.axis('off')

    plt.tight_layout()
    plt.show()

# Load a batch of images and labels from the training dataset
sample_loader = iter(train_loader)
images, labels = next(sample_loader)

# Get class names from the dataset
class_names = train_dataset.classes

# Visualize the images
imshow_grid(images, labels, class_names)
```

Crazing	Crazing	Rolled	Crazing
Scratches	Rolled	Patches	Crazing
Patches	Crazing	Patches	Inclusion
Patches	Rolled	Rolled	Scratches

```python
# Using VGG16
model_vgg16 = vgg16(weights=True)
num_classes = len(train_dataset.classes)

# Replace the last fully connected layer to match your dataset's
# number of classes
model_vgg16.classifier[6] = nn.Linear(4096, num_classes)

device = torch.device("cuda:0" if torch.cuda.is_available() else
"cpu")
model_vgg16 = model_vgg16.to(device)
```

```python
# Define Loss Function and optimizer

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(model_vgg16.parameters(), lr=0.001,
momentum=0.9)

# Train model

def train_model(model, criterion, optimizer, num_epochs=10):
    train_loss_list = []
    valid_loss_list = []
    train_acc_list = []
    valid_acc_list = []

    for epoch in range(num_epochs):
        model.train()
        running_loss = 0.0
        correct_train = 0
        total_train = 0

        for inputs, labels in train_loader:
            inputs, labels = inputs.to(device), labels.to(device)

            optimizer.zero_grad()
            outputs = model(inputs)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()

            running_loss += loss.item()

            _, predicted = torch.max(outputs, 1)
            total_train += labels.size(0)
            correct_train += (predicted == labels).sum().item()

        epoch_loss = running_loss / len(train_loader)
        train_accuracy = correct_train / total_train

        # Validation
        model.eval()
        correct_valid = 0
        total_valid = 0
        running_valid_loss = 0.0

        with torch.no_grad():
            for inputs, labels in valid_loader:
                inputs, labels = inputs.to(device), labels.to(device)

                outputs = model(inputs)
                loss = criterion(outputs, labels)
                running_valid_loss += loss.item()
```

```python
                _, predicted = torch.max(outputs, 1)
                total_valid += labels.size(0)
                correct_valid += (predicted == labels).sum().item()

        epoch_valid_loss = running_valid_loss / len(valid_loader)
        valid_accuracy = correct_valid / total_valid

        # Record metrics
        train_loss_list.append(epoch_loss)
        valid_loss_list.append(epoch_valid_loss)
        train_acc_list.append(train_accuracy)
        valid_acc_list.append(valid_accuracy)

        print(f"Epoch {epoch + 1}/{num_epochs}, Train Loss:
{epoch_loss:.4f}, Valid Loss: {epoch_valid_loss:.4f}, Train Accuracy:
{train_accuracy:.4f}, Valid Accuracy: {valid_accuracy:.4f}")

    # Plot loss and accuracy curves
    plt.figure(figsize=(12, 5))
    plt.subplot(1, 2, 1)
    plt.plot(range(1, num_epochs + 1), train_loss_list, label='Train')
    plt.plot(range(1, num_epochs + 1), valid_loss_list,
label='Validation')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.legend()
    plt.title('Loss Curve')

    plt.subplot(1, 2, 2)
    plt.plot(range(1, num_epochs + 1), train_acc_list, label='Train')
    plt.plot(range(1, num_epochs + 1), valid_acc_list,
label='Validation')
    plt.xlabel('Epoch')
    plt.ylabel('Accuracy')
    plt.legend()
    plt.title('Accuracy Curve')

    plt.tight_layout()
    plt.show()

train_model(model_vgg16, criterion, optimizer, num_epochs=10)

Epoch 1/10, Train Loss: 0.7899, Valid Loss: 0.0010, Train Accuracy:
0.7560, Valid Accuracy: 1.0000
Epoch 2/10, Train Loss: 0.0545, Valid Loss: 0.0009, Train Accuracy:
0.9795, Valid Accuracy: 1.0000
Epoch 3/10, Train Loss: 0.0141, Valid Loss: 0.0000, Train Accuracy:
0.9952, Valid Accuracy: 1.0000
Epoch 4/10, Train Loss: 0.0181, Valid Loss: 0.0001, Train Accuracy:
```

```
0.9958, Valid Accuracy: 1.0000
Epoch 5/10, Train Loss: 0.0021, Valid Loss: 0.0000, Train Accuracy:
1.0000, Valid Accuracy: 1.0000
Epoch 6/10, Train Loss: 0.0011, Valid Loss: 0.0000, Train Accuracy:
1.0000, Valid Accuracy: 1.0000
Epoch 7/10, Train Loss: 0.0005, Valid Loss: 0.0000, Train Accuracy:
1.0000, Valid Accuracy: 1.0000
Epoch 8/10, Train Loss: 0.0003, Valid Loss: 0.0000, Train Accuracy:
1.0000, Valid Accuracy: 1.0000
Epoch 9/10, Train Loss: 0.0006, Valid Loss: 0.0000, Train Accuracy:
1.0000, Valid Accuracy: 1.0000
Epoch 10/10, Train Loss: 0.0003, Valid Loss: 0.0000, Train Accuracy:
1.0000, Valid Accuracy: 1.0000
```



```python
# Evaluate the model

def evaluate_model(model, data_loader):
    model.eval()
    correct_predictions = 0
    total_samples = 0
    all_predictions = []
    all_labels = []

    with torch.no_grad():
        for inputs, labels in data_loader:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, predicted = torch.max(outputs, 1)
            total_samples += labels.size(0)
            correct_predictions += (predicted == labels).sum().item()

            all_predictions.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())
```

```python
    accuracy = correct_predictions / total_samples
    precision = precision_score(all_labels, all_predictions,
average='weighted')
    recall = recall_score(all_labels, all_predictions,
average='weighted')
    f1 = f1_score(all_labels, all_predictions, average='weighted')

    print(f"Accuracy: {accuracy:.4f}")
    print(f"Precision: {precision:.4f}")
    print(f"Recall: {recall:.4f}")
    print(f"F1 Score: {f1:.4f}")

# Evaluate on test set
evaluate_model(model_vgg16, test_loader)

Accuracy: 1.0000
Precision: 1.0000
Recall: 1.0000
F1 Score: 1.0000
```